


# Kubernetes & Golang

A retrospective

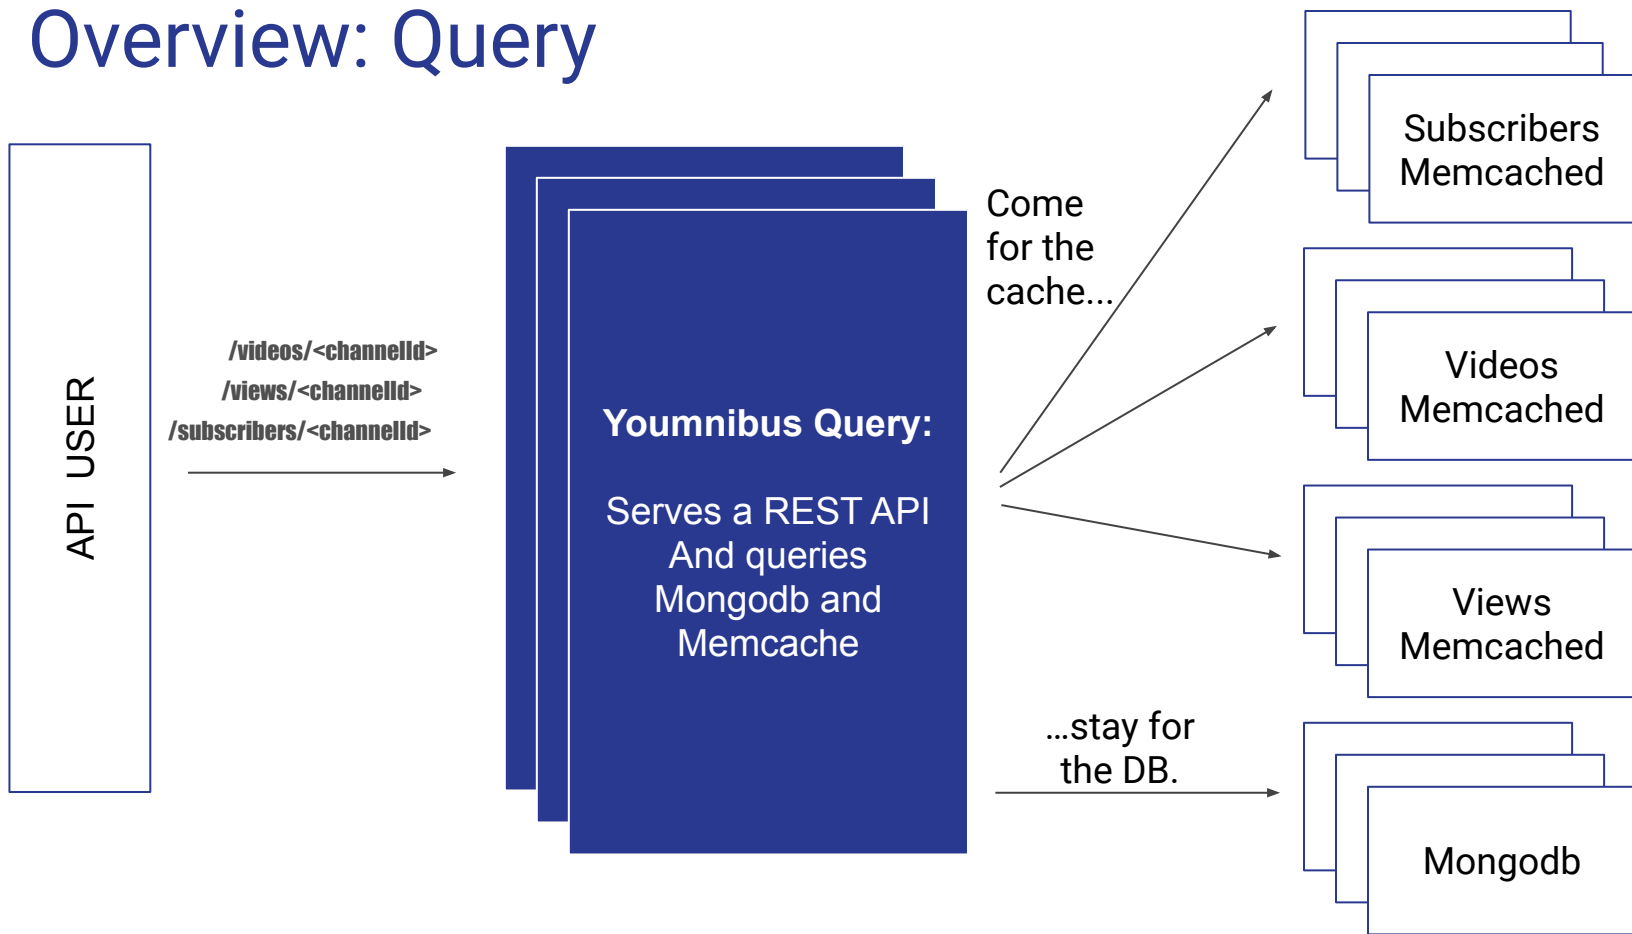
# My project

Youmnibus = YouTube + Omnibus. It is a system for capturing and querying time series YouTube data.

The stack:

- YouTube Data API
  - Google Kubernetes Engine
  - Mongoddb
  - RabbitMQ
  - Golang
  - Memcached
- 

# Overview: Query



# Overview: Query

## Quirks:

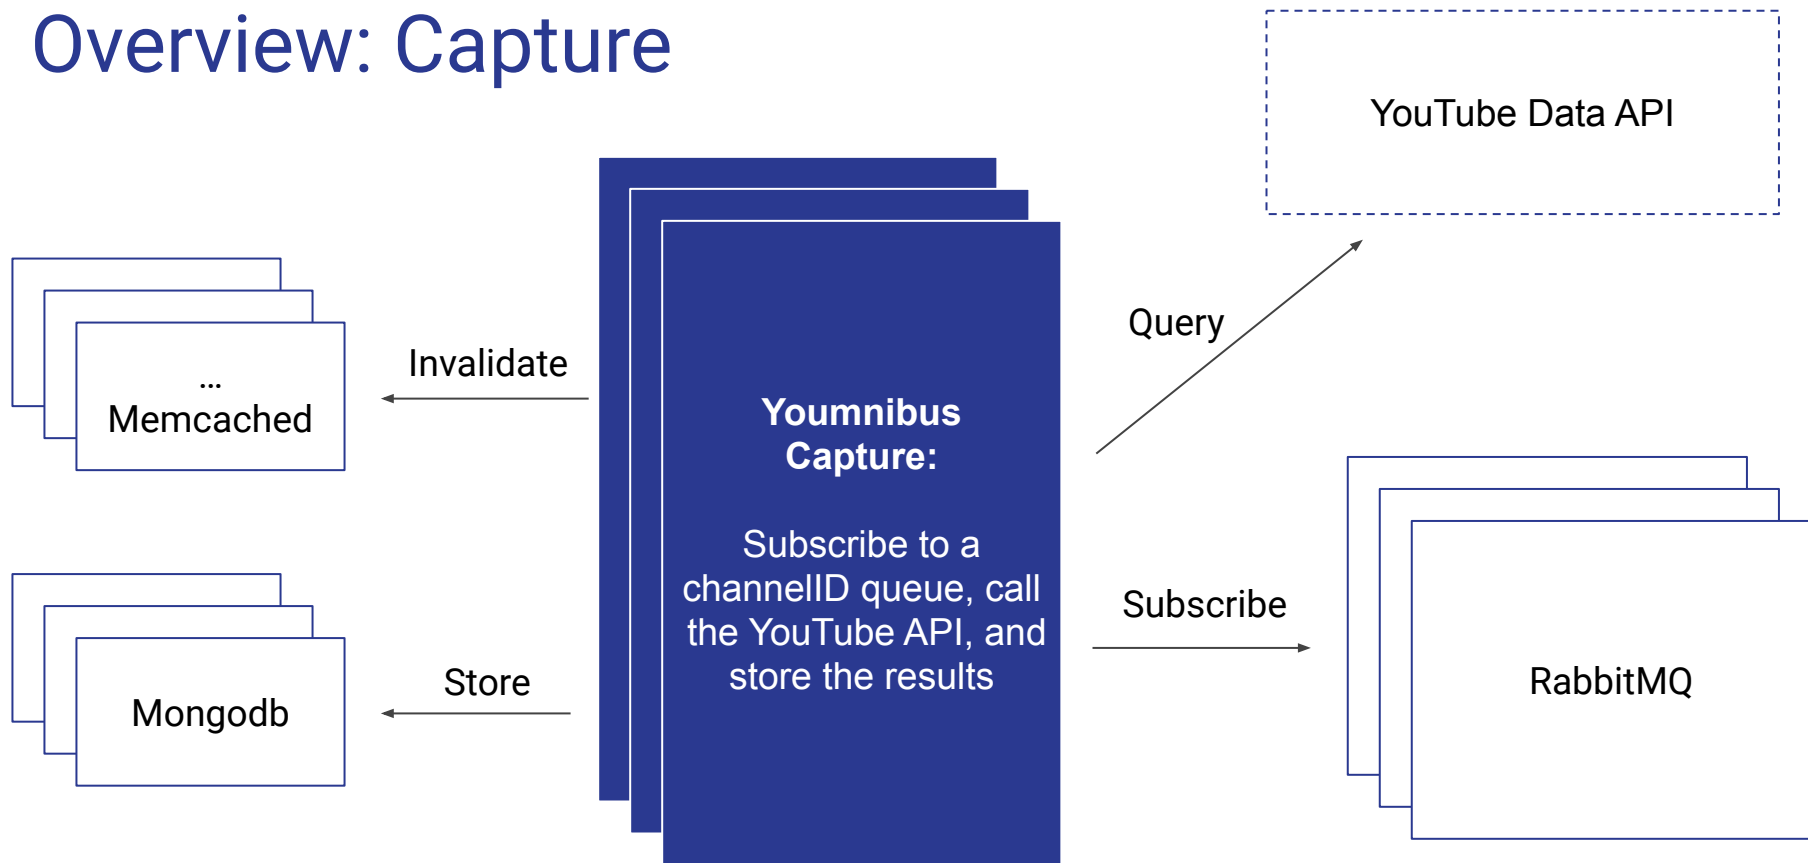
- JSON serialization with base Golang kinda sucks

```
type SubscribersAt struct {  
    SubscriberCount uint64 `json:"subscriberCount,omitempty,string"`  
    At              string  `json:"at,omitempty"`  
}
```

- I use Gin as a REST framework. It's meant to be a relatively lightweight and high performance router. You can plug in middleware for e.g. logging and authentication - I've just used the defaults.




# Overview: Capture

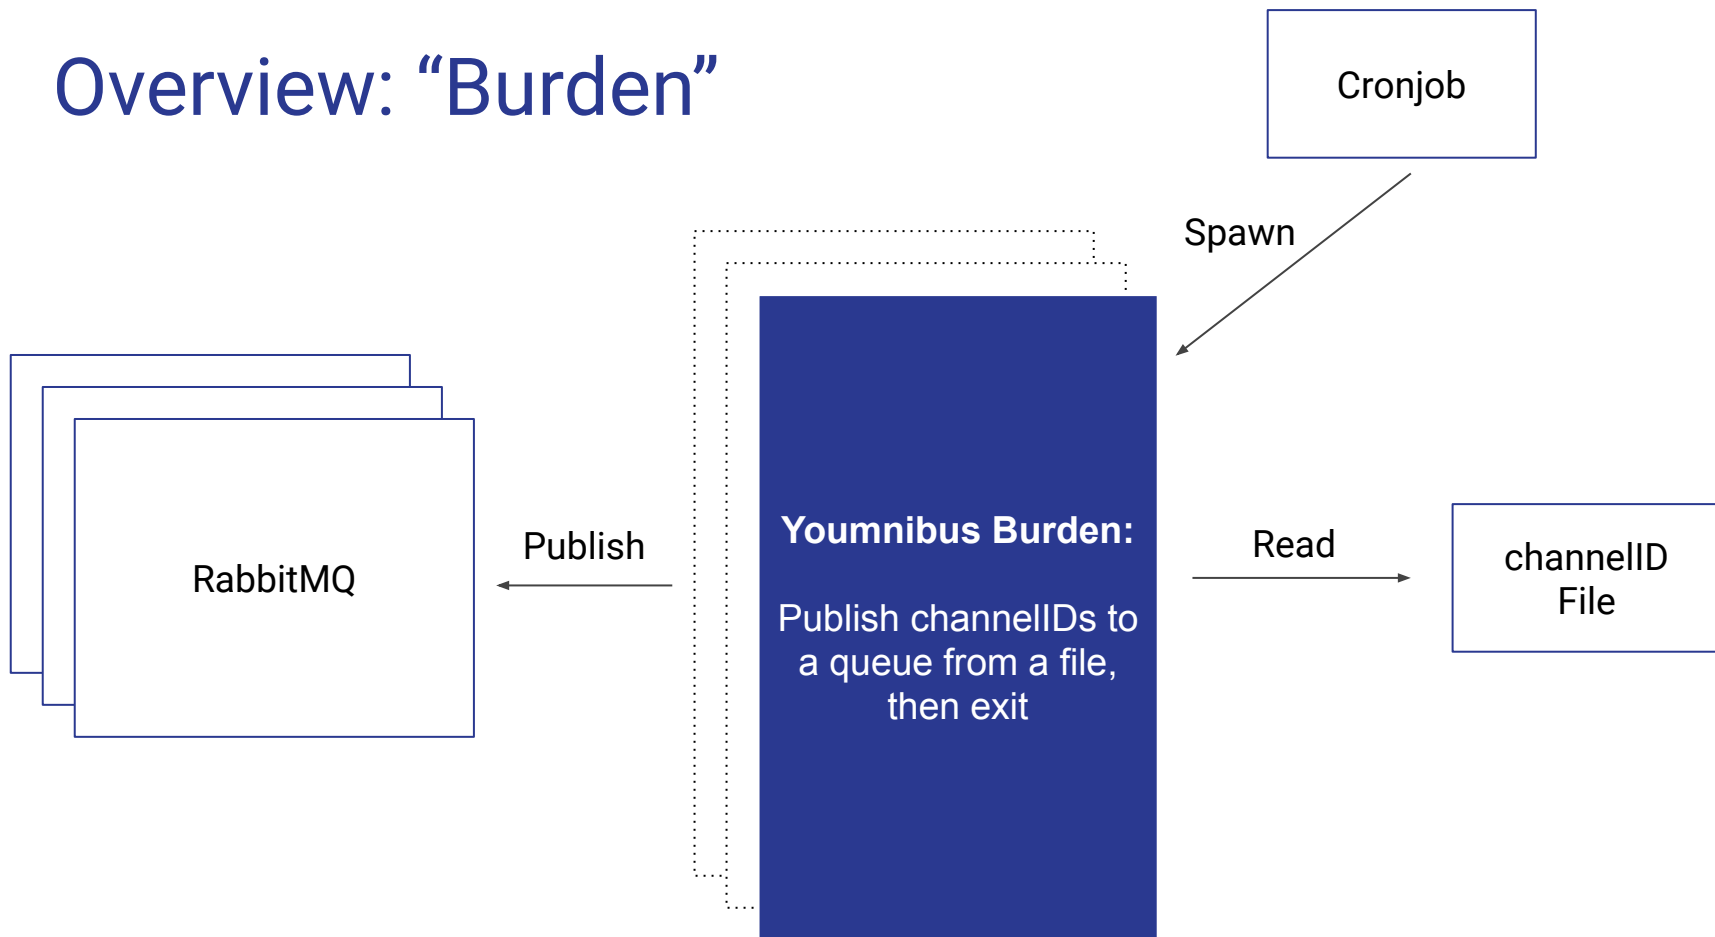


# Overview: Capture

## Quirks:

- I authenticate and authorize with the YouTube API using an API key. If your system is querying/updating more personal user info, then there is an Oauth option.
  - You can create all the RabbitMQ needed resources from the Golang client itself. I.e. A queue and a dead letter exchange. The first application to request the queue connection will construct it, the rest will just use it.
  - Golang concurrency + RabbitMQ worker queues work really nicely - Golang subscribes to a (Golang) channel which is populated asynchronously. It's reliable and the channel doesn't use any resources while idle.
- 

# Overview: “Burden”



# Overview: Burden

## Quirks:

- Because Golang applications are statically linked (to C libs) they work great for worker applications: Golang apps start up quickly and use little base memory.
  - You can also compile *with* the C libs, so that the binary has no dependencies whatsoever.





# Environment Evolution



Local

**“I’m too young to die”**

Get the applications running in separate terminals.

Docker

**“Hey, not too rough”**

Get the applications running in Docker, and create a build pipeline using Docker Hub and GitHub.

Kubernetes

**“Hurt me plenty”**

Deploy load balanced applications on Google Kubernetes Engine with volume claims, secrets and config maps.

# Local

## Quirks:

- I configure my applications using environment variables which have as many common sense defaults as possible, e.g. `MONGO_PORT: 27017`. This makes it easy to run the application locally - by just running the executable - but also easy to configure the app in both Docker and Kubernetes.
- I find it is helpful to log the config at the start of the application, it serves as a kind of minimal manual about how it can be configured.



# Docker

## Quirks:

- A Makefile is a great way of automating common tasks for your application, including:
  - Building and installing the application locally
  - Building / Deploying / Pushing docker images
  - ...and anything else which can be run in a few lines in the terminal.
  - Plus it has bash completion



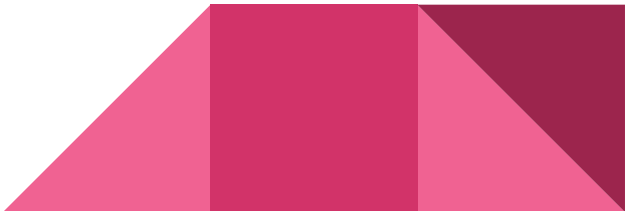
# Docker

## Quirks:

- E.g. make install docker-build docker-push

Is equivalent to:

```
go build ./...  
go install ./...  
docker build -t youmnibus-burden -f docker/Dockerfile.youmnibus-burden .  
docker build -t youmnibus-capture -f docker/Dockerfile.youmnibus-capture .  
docker build -t youmnibus-query -f docker/Dockerfile.youmnibus-query .  
docker tag youmnibus-burden:latest lpulles/youmnibus-burden:0.1  
docker tag youmnibus-capture:latest lpulles/youmnibus-capture:0.1  
docker tag youmnibus-query:latest lpulles/youmnibus-query:0.1  
docker push lpulles/youmnibus-burden:0.1  
docker push lpulles/youmnibus-capture:0.1  
docker push lpulles/youmnibus-query:0.1
```



# Docker

Here's a nice Dockerfile for Golang apps:

```
1 FROM golang:alpine
2 RUN apk update && apk add --no-cache git
3
4 ENV GOPATH /go
5 ENV GOBIN /go/bin
6 ENV CGO_ENABLED 0
7 WORKDIR /app
8 RUN mkdir -p /go/bin
9
10 COPY ./go.mod .
11 RUN go mod download
12 RUN go mod verify
13
14 COPY . .
15 RUN go build ./...
16 RUN go install ./...
17
18 FROM scratch
19 COPY --from=0 /go/bin/youmnibus-burden /go/bin/youmnibus-burden
20 ENTRYPOINT ["/go/bin/youmnibus-burden"]
```

# Docker

## Quirks:

- Some aspects of the minimal image to consider
  - The resulting image is about ~25MB large
  - You can't run a terminal inside the container, so you can't list files or check what environment variables the container is currently using. This makes it difficult to debug, though also harder to attack.
  - The benefits/pitfalls of this approach don't really relate to the size of the image, because if I instead used e.g. an ubuntu image, I would only download the application layer of the image, and would only hold in memory the application layer while running the container.
  - So, do this at your discretion.



# Kubernetes

## Quirks:

- Constructing secrets for strings (e.g. passwords) can be deceptively tricky.
  - To construct a secret for a password, you need to provide a base64 encoded version of the string.
  - For those with some bash experience, this might seem trivial e.g. we just run `echo "duck" | base64`
  - The trick is, the echo command inserts a newline character at the end of the string, which changes the actual password.
  - You need:  
`echo -n "duck" | base64`
  - This caused many hours of frustration.



# Kubernetes

## Quirks:

- Connecting to the instances of a stateful set can be deceptively tricky.
  - A stateful set is a deployment of an application where the instances aren't load balanced and are not automatically killed for new deployments. E.g. a mongodb deployment.
  - For mongodb, you can use a connection string in your apps like `mongo://user:pass@host1,host2,host3` where the client will select a host accordingly. So what DNS address do we give the application for host1, host2 and host3?
  - It's NOT `mongod-0`, `mongod-1` and `mongod-2` i.e. the names of the individual pods
  - It's NOT given by an environment variable - unlike normal services which are.
  - What it actually is, is `mongod-0.mongodb-service.yournibus.svc.cluster.local`, `mongod-1.mongodb-service.yournibus.svc.cluster.local`, `mongod-2.mongodb-service.yournibus.svc.cluster.local` (where `yournibus` is the kubernetes namespace I'm using).
  - This caused many hours of frustration.



# Kubernetes

## Quirks:

- The initial stages of setting up Kubernetes is in general deceptively tricky and causes many hours of frustration.
- BUT, when it finally works, it's awesome!



# Demo

<https://github.com/liampulles/yournibus>